

Machine Learning Course - CS-433

Text Representation Learning

Nov 24, 2016

©Martin Jaggi 2016



ÉCOLE POLYTECHNIQUE
FÉDÉRALE DE LAUSANNE


Motivation

Finding numerical representations for words is fundamental for all machine learning methods dealing with text data.

Goal: For each word, find mapping (embedding)

$$w_i \mapsto \mathbf{w}_i \in \mathbb{R}^K$$

Representation should capture semantics of the word.



Constructing good feature representations (= [representation learning](#)) benefits all ML applications.

The Co-Occurrence Matrix

A big corpus of un-labeled text can be represented as the [co-occurrence counts](#)

$n_{ij} := \# \text{contexts where word } w_i \text{ occurs together with word } w_j.$

	1	1		
		3		
	1			
	2		1	
1				1
		1		
	1		1	1

Needs definition of

- [Context](#) e.g. document, paragraph, sentence, window
- [Vocabulary](#)

$$\mathcal{V} := \{w_1, \dots, w_D\}$$

For [words](#) $w_d = 1, 2, \dots, D$ and [context words](#) $w_n = 1, 2, \dots, N$, the co-occurrence counts n_{ij} form a very sparse $D \times N$ matrix.

Learning Word-Representations (Using Matrix Factorization)

Find a factorization of the co-occurrence matrix!

Typically uses log of the actual counts, i.e. $x_{dn} := \log(n_{dn})$.

We will aim to find \mathbf{W}, \mathbf{Z} s.t.

$$\mathbf{X} \approx \mathbf{W}\mathbf{Z}^\top.$$

So for each pair of words (w_d, w_n) , we try to ‘explain’ their co-occurrence count by a numerical representation of the two words

- in fact by the inner product of the two feature vectors $\mathbf{W}_{d:}, \mathbf{Z}_{n:}$.

$$\min_{\mathbf{W}, \mathbf{Z}} \mathcal{L}(\mathbf{W}, \mathbf{Z}) := \frac{1}{2} \sum_{(d,n) \in \Omega} f_{dn} [x_{dn} - (\mathbf{W}\mathbf{Z}^\top)_{dn}]^2$$

where $\mathbf{W} \in \mathbb{R}^{D \times K}$ and $\mathbf{Z} \in \mathbb{R}^{N \times K}$ are tall matrices, having only $K \ll D, N$ columns.

The set $\Omega \subseteq [D] \times [N]$ collects the indices of non-zeros of the count matrix \mathbf{X} .

Each row of those matrices forms a [representation of a word](#) (\mathbf{W}) or a context word (\mathbf{Z}) respectively.


GloVe

This model is called **GloVe**, and is a variant of **word2vec**.

Weights f_{dn} : Give “importance” of each entry. Choosing $f_{dn} := 1$ is ok.

GloVe weight function:


$$f_{dn} := \min \{1, (n_{dn}/n_{\max})^\alpha\}, \quad \alpha \in [0; 1] \quad \text{e.g. } \alpha = \frac{3}{4}$$




Choosing K

K e.g. 50, 100, 200


Word Analogies



Male-Female



Verb tense



Country-Capital

Newspapers			
New York	New York Times	Baltimore	Baltimore Sun
San Jose	San Jose Mercury News	Cincinnati	Cincinnati Enquirer
NHL Teams			
Boston	Boston Bruins	Montreal	Montreal Canadiens
Phoenix	Phoenix Coyotes	Nashville	Nashville Predators
NBA Teams			
Detroit	Detroit Pistons	Toronto	Toronto Raptors
Oakland	Golden State Warriors	Memphis	Memphis Grizzlies
Airlines			
Austria	Austrian Airlines	Spain	Spainair
Belgium	Brussels Airlines	Greece	Aegean Airlines
Company executives			
Steve Ballmer	Microsoft	Larry Page	Google
Samuel J. Palmisano	IBM	Werner Vogels	Amazon

Training

- Stochastic Gradient Descent (SGD)
- Alternating Least-Squares (ALS)

Open questions:

- Parallel and distributed training
- Does regularization help?

Alternative: Skip-Gram Model

(Original word2vec)

Uses binary classification (logistic regression objective), to separate good word pairs (w_d, w_n) from bad word pairs. Same inner product score = matrix factorization.


Given w_d , a context word w_n is

- Good = appearing together in a context window of 5
- Bad = any word $w_{n'}$ sampled randomly: [Negative sampling](#)

Learning Representations of Documents

Supervised:

For a supervised task (e.g. predicting the emotion of a tweet), we can use matrix-factorization (below) or convolutional neural networks (see next weeks).



→ SemEval competition for tweet classification.

Unsupervised:

Open research.

FastText

Matrix factorization to learn document/sentence representations (supervised).

Given a sentence $s_n = (w_1, w_2, \dots, w_m)$, let $\mathbf{x}_n \in \mathbb{R}^{|\mathcal{V}|}$ be the bag-of-words representation of the sentence.

$$\min_{\mathbf{W}, \mathbf{Z}} \mathcal{L}(\mathbf{W}, \mathbf{Z}) := \sum_{s_n \text{ a sentence}} f(y_n \mathbf{W} \mathbf{Z}^\top \mathbf{x}_n)$$

where $\mathbf{W} \in \mathbb{R}^{1 \times K}$, $\mathbf{Z} \in \mathbb{R}^{|\mathcal{V}| \times K}$ are the variables, and the vector $\mathbf{x}_n \in \mathbb{R}^{|\mathcal{V}|}$ represents our n -th training sentence.

Here f is a linear classifier loss function, and y_n is the classification label for sentence \mathbf{x}_n .

Further Pointers

1. word2vec:

code: code.google.com/p/word2vec/

paper:

“Distributed representations of words and phrases and their compositionality” - T Mikolov, I Sutskever, K Chen, GS Corrado, J Dean. NIPS 2013

2. GloVe:

code and vectors: nlp.stanford.edu/projects/glove/

paper:

“GloVe: Global Vectors for Word Representation” - Pennington, J., Socher, R., Manning, C. D.. EMNLP 2014

3. FastText

code: github.com/facebookresearch/fastText

papers:

“Bag of Tricks for Efficient Text Classification” - Joulin, A., Grave, E., Bojanowski, P., Mikolov, T. - arXiv, 2016.

“Enriching Word Vectors with Subword Information” - Bojanowski, P., Grave, E., Joulin, A., Mikolov, T. - arXiv, 2016.