

*Annotated
Version*

Machine Learning Course - CS-433

Text Representation Learning

Nov 23, 2017

©Martin Jaggi 2016

minor changes by Martin Jaggi 2017

Last updated: November 23, 2017



ÉCOLE POLYTECHNIQUE
FÉDÉRALE DE LAUSANNE

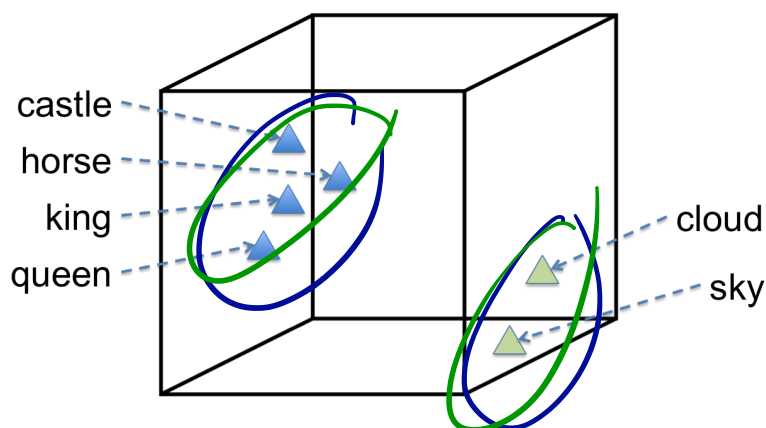

Motivation

Finding numerical representations for words is fundamental for all machine learning methods dealing with text data.

Goal: For each word, find mapping (embedding)

$$w_i \mapsto \underline{\mathbf{w}_i} \in \mathbb{R}^{\overset{K}{\text{K}}}$$

Representation should capture semantics of the word.



Constructing good feature representations (= representation learning) benefits all ML applications.

The Co-Occurrence Matrix


A big corpus of un-labeled text can be represented as the **co-occurrence counts**

$n_{ij} := \# \text{contexts where word } w_i \text{ occurs together with word } w_j.$

w_j (context words)

0	1	1	0	0
0	0	3	0	0
	1			
	2		1	
1				1
		1		
	1		1	1

w_i →



Needs definition of

- **Context** e.g. document, paragraph, sentence, window

- **Vocabulary**

$$\mathcal{V} := \{w_1, \dots, w_D\}$$

typical: window size = 5

For words $w_d = 1, 2, \dots, D$ and context words $w_n = 1, 2, \dots, N$, the co-occurrence counts n_{ij} form a very sparse $D \times N$ matrix.

$$N = D = |\mathcal{V}|$$

Learning Word-Representations (Using Matrix Factorization)

Find a factorization of the co-occurrence matrix!


Typically uses log of the actual counts, i.e. $x_{dn} := \log(n_{dn})$.

We will aim to find \mathbf{W}, \mathbf{Z} s.t.

$$\mathbf{X} \approx \mathbf{W}\mathbf{Z}^\top$$

So for each pair of words (w_d, w_n) , we try to 'explain' their co-occurrence count by a numerical representation of the two words

- in fact by the inner product of the two feature vectors $\mathbf{W}_{d:}, \mathbf{Z}_{n:}$.



$$\min_{\mathbf{W}, \mathbf{Z}} \mathcal{L}(\mathbf{W}, \mathbf{Z}) := \frac{1}{2} \sum_{(d,n) \in \Omega} f_{dn} [x_{dn} - (\mathbf{W}\mathbf{Z}^\top)_{dn}]^2$$

where $\mathbf{W} \in \mathbb{R}^{D \times K}$ and $\mathbf{Z} \in \mathbb{R}^{N \times K}$ are tall matrices, having only $K \ll D, N$ columns.

The set $\Omega \subseteq [D] \times [N]$ collects the indices of non-zeros of the count matrix \mathbf{X} .

Each row of those matrices forms a representation of a word (\mathbf{W}) or a context word (\mathbf{Z}) respectively.

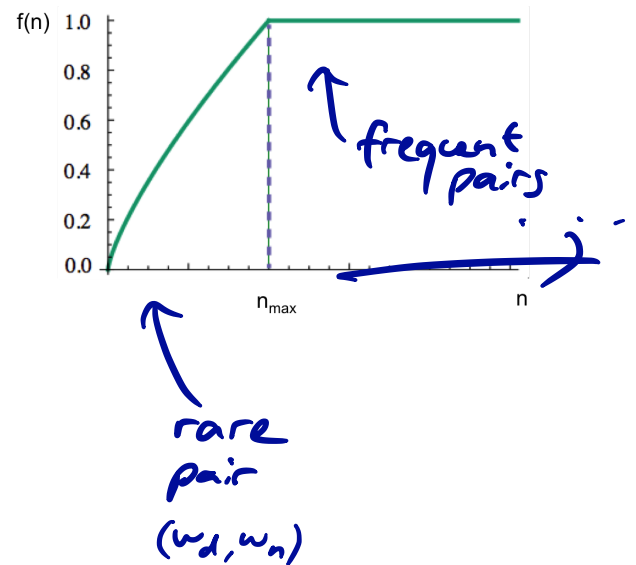
GloVe

This model is called GloVe, and is a variant of word2vec.

Weights f_{dn} : Give “importance” of each entry. Choosing $f_{dn} := 1$ is ok.
GloVe weight function:

$$f_{dn} := \min \{1, (n_{dn}/n_{\max})^\alpha\}, \quad \alpha \in [0; 1] \quad \text{e.g. } \alpha = \frac{3}{4}$$


count
(w_d, w_n)




Choosing K

K e.g. 50, 100, 200


Word Analogies



Male-Female




Verb tense



Country-Capital

Newspapers			
New York San Jose	New York Times San Jose Mercury News	Baltimore Cincinnati	Baltimore Sun Cincinnati Enquirer
NHL Teams			
Boston Phoenix	Boston Bruins Phoenix Coyotes	Montreal Nashville	Montreal Canadiens Nashville Predators
NBA Teams			
Detroit Oakland	Detroit Pistons Golden State Warriors	Toronto Memphis	Toronto Raptors Memphis Grizzlies
Airlines			
<u>Austria</u> Belgium	<u>Austrian Airlines</u> Brussels Airlines	Spain Greece	<u>Spanair</u> Aegean Airlines
Company executives			
Steve Ballmer Samuel J. Palmisano	Microsoft IBM	Larry Page Werner Vogels	Google Amazon



Training

- Stochastic Gradient Descent (SGD)
- Alternating Least-Squares (ALS)

same as for
Recommender
System

Open questions:

- Parallel and distributed training
- Does regularization help?

Alternative: Skip-Gram Model


(Original word2vec)

(CBOW)

Uses binary classification (logistic regression objective), to separate **real** word pairs (w_d, w_n) from **fake** word pairs. Same inner product score = matrix factorization.

Given w_d , a context word w_n is


- **real** = appearing together in a context window of size 5
- **fake** = any word $w_{n'}$ sampled randomly: Negative sampling (also: Noise Contrastive Estimation)



Learning Representations of Documents

Supervised: For a supervised task (e.g. predicting the emotion of a tweet), we can use matrix-factorization (below) or convolutional neural networks (see next weeks).


CNN



→ SemEval competition for tweet classification.

Unsupervised:

- Adding (fixed, given) word vectors over sentence or document
- Training word vectors such that adding works well
- Direct unsupervised training for sentences (appearing together with context sentences) instead of words

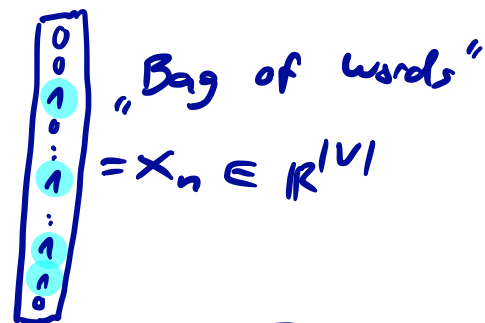


FastText


Matrix factorization to learn document/sentence representations (supervised).

Given a sentence $s_n = (w_1, w_2, \dots, w_m)$, let $\mathbf{x}_n \in \mathbb{R}^{|\mathcal{V}|}$ be the bag-of-words representation of the sentence.


supervised
(document classification)



$$\min_{\mathbf{W}, \mathbf{Z}} \mathcal{L}(\mathbf{W}, \mathbf{Z}) := \sum_{s_n \text{ a sentence}} f(y_n \mathbf{W} \mathbf{Z}^T \mathbf{x}_n)$$




where $\mathbf{W} \in \mathbb{R}^{1 \times K}$, $\mathbf{Z} \in \mathbb{R}^{|\mathcal{V}| \times K}$ are the variables, and the vector $\mathbf{x}_n \in \mathbb{R}^{|\mathcal{V}|}$ represents our n -th training sentence.



SGD

Here f is a linear classifier loss function, and $y_n \in \{\pm 1\}$ is the classification label for sentence \mathbf{x}_n .

like in
Logistic
regression



Further Pointers

1. word2vec:

code: code.google.com/p/word2vec/

paper:

“Distributed representations of words and phrases and their compositionality” - T Mikolov, I Sutskever, K Chen, GS Corrado, J Dean. NIPS 2013

2. GloVe:

code and vectors: nlp.stanford.edu/projects/glove/

paper:

“GloVe: Global Vectors for Word Representation” - Pennington, J., Socher, R., Manning, C. D.. EMNLP 2014

3. FastText

code: github.com/facebookresearch/fastText

papers:

“Bag of Tricks for Efficient Text Classification” - Joulin, A., Grave, E., Bojanowski, P., Mikolov, T. - [EC-ACL](#), 2017.

“Enriching Word Vectors with Subword Information” - Bojanowski, P., Grave, E., Joulin, A., Mikolov, T. - [TACL](#), 2017.

“Unsupervised Learning of Sentence Embeddings using Compositional n-Gram Features” - Pagliardini, M., Gupta, P., Jaggi, M. [arXiv](#) 2017.